

# CS2030 Lecture 10

## Parallel and Concurrent Programming

Henry Chia (hchia@comp.nus.edu.sg)

Semester 1 2018 / 2019

## Concurrency vs Parallelism

- A single core processor executes one instruction at a time
  - Only one process can run at any one time
  - Context-switching allows multi-tasking on a single processor
- Concurrent programs run concurrently via threads
  - OS switches between threads
  - Separate unrelated tasks into separate threads
  - Improves processor utilization
- Parallel computing involves multiple subtasks running at the same time on multiple (possibly multi-core) processors
- Parallel programs are concurrent, but not all concurrent programs are parallel


1 / 24

3 / 24

## Lecture Outline

- Concurrency versus parallelism
- Parallel streams
- Fork/join framework
- Thread pools

## Concurrency vs Parallelism



2 / 24

4 / 24

## Parallel Streams

- Streams can be executed in parallel to increase runtime performance
- Parallel streams use a common `ForkJoinPool` via the static `ForkJoinPool.commonPool()` method

```
ForkJoinPool commonPool = ForkJoinPool.commonPool();
System.out.println(commonPool.getParallelism());
```
- Collections support the method `parallelStream()` to create a parallel stream of elements
- Alternatively, the intermediate operation `parallel` can be invoked on a given stream to parallelize a sequential stream

5 / 24

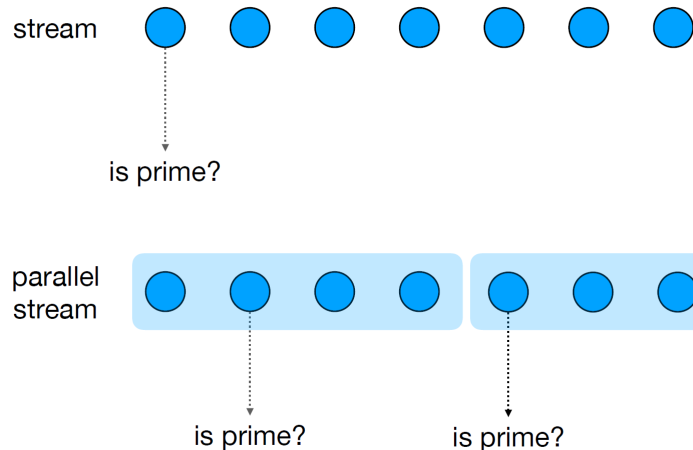
## Parallel Streams

```
int sum = IntStream.of(1, 2, 3, 4, 5)
    .parallel()
    .filter(x -> {
        System.out.println("filter: " + x + " "
            + Thread.currentThread().getName());
        return true;
    })
    .map(x -> {
        System.out.println("map: " + x + " "
            + Thread.currentThread().getName());
        return x;
    })
    .reduce(9, (x, y) -> {
        System.out.println("reduce: " + x + " + " + y + " "
            + Thread.currentThread().getName());
        return x + y;
    });
System.out.println(sum);
```

7 / 24

## Parallel Streams

- Using prime number testing as an example:



6 / 24

## Correctness of Parallel Streams

- To ensure correct parallel execution, stream operations must not interfere with stream data, preferably stateless and have no side effects
- Example:

```
List<String> list = new ArrayList<>(
    List.of("abc", "def", "xyz"));

list.stream()
    .peek(str -> {
        if (str.equals("xyz")) {
            list.add("pqr");
        }
    })
    .forEach(x -> {});
```
- Inteferece is not allowed in both sequential and parallel streams

8 / 24

## Correctness of Parallel Streams

- Another example:

```
List<Integer> list = new ArrayList<>(  
    Arrays.asList(1, 3, 5, 7, 9, 11, 13, 15, 17, 19));  
List<Integer> result = new ArrayList<>();
```
- The following is erroneous

```
list.parallelStream() // list.stream().parallel()  
    .filter(x -> isPrime(x))  
    .forEach(x -> result.add(x));
```
- Use `.collect` instead

```
result = list.parallelStream()  
    .filter(x -> isPrime(x))  
    .collect(Collectors.toList());
```
- Side effects are a problem in parallel streams
- Use a thread-safe list (e.g. `CopyOnWriteArrayList`)

9 / 24

## Accumulator and Combiner

- Accumulator and combiner functions are executed in parallel

```
int result = Stream.of(1, 2, 3, 4)  
    .parallel()  
    .reduce(  
        1,  
        (x,y) -> {  
            System.out.println("accumulator: " +  
                x + " * " + y);  
            return x * y;  
        },  
        (x,y) -> {  
            System.out.println("combiner: " +  
                x + " * " + y);  
            return x * y;  
        });  
System.out.println(result);
```

11 / 24

## Inherently Parallelizable reduce


- Consider Stream's three-argument reduce method:

```
<U> U reduce(U identity,  
    BiFunction<U,? super T,U> accumulator,  
    BinaryOperator<U> combiner)
```
- Rules to follow when parallelizing
  - `combiner.apply(identity, i)` must be equal to `i`
  - combiner and accumulator must be associative, i.e. order of application does not matter
  - combiner and accumulator must be compatible, i.e. `combiner.apply(u, accumulator.apply(identity, t))` must be equal to `accumulator.apply(u, t)`
  - The following example compiles with the above rules:

```
Stream.of(1,2,3,4)  
    .parallel()  
    .reduce(1, (x,y) -> x * y, (x,y) -> x * y)
```

10 / 24

## Accumulator and Combiner



12 / 24


## Accumulator and Combiner

- Errneous examples:

```
double result = Stream.of(1, 2, 3, 4)
    .reduce(1.0,
        (x, y) -> x + y,
        (x, y) -> x + y);
```

```
result = Stream.of(1, 2, 3, 4)
    .parallel()
    .reduce(24.0,
        (x, y) -> 1.0 * x / y,
        (x, y) -> 1.0 * x / y);
```

## Thread Pools




13 / 24

15 / 24

## Fork and Join

- Given the following program fragment and *computation* graph



- $f(a)$  invoked before  $g(b)$  and  $h(b)$ ;  $n(c,d)$  invoked after
- How about the order of  $g(b)$  and  $h(b)$ ?
  - If  $g$  and  $h$  does not produce side effects, then parallelize
  - **Fork** task  $g$  to execute at the same time as  $h$ , and **join** back task  $g$  later

## Thread Pools

- Java maintains a pool of *worker threads*
  - Each thread is an abstraction of a running task
  - Task submitted to the pool for execution, and joins a queue (global queue or worker queue)
  - Worker thread picks a task from the queue to execute
- ForkJoinPool is the class that implements the thread pool for RecursiveTask (a sub-class of ForkJoinTask)
- To submit a task to the thread pool:

```
int sum = ForkJoinPool.commonPool().invoke(task);
```
- `invoke(task)` versus `task.compute()`
  - `task.compute()` invokes task immediately; may result in stack overflow if too many recursive tasks
  - `invoke(task)` gets the task to join the queue, waiting to be carried out by a worker (recommended)

14 / 24

16 / 24

## Fork/Join Framework

```
import java.util.concurrent.RecursiveTask;

class SumLeftRight extends RecursiveTask<Integer> {
    int low;
    int high;
    int[] array;

    SumLeftRight(int low, int high, int[] array) {
        this.low = low;
        this.high = high;
        this.array = array;
    }

    @Override
    protected Integer compute() {
        if (high - low < 2) {
            int sum = 0;
            for (int i = low; i < high; i++) {
                sum += array[i];
            }
            return sum;
        } else {
            int middle = (low + high) / 2;
            SumLeftRight left = new SumLeftRight(low, middle, array);
            SumLeftRight right = new SumLeftRight(middle, high, array);
            left.fork();
            return right.compute() + left.join();
        }
    }
}
```

17 / 24


## Fork/Join in Parallel Streams

- `parallel()` runs fork to create sub-tasks running the same chain of operations on sub-streams
  - combiner in reduce runs join to combine the results
  - Parallelizing a trivial task actually creates more work in terms of parallelizing overhead
- ```
IntStream.range(2, (int) Math.sqrt(n) + 1)
    .parallel()
    .noneMatch(x -> n % x == 0);
```
- Parallelization is worthwhile if the task is complex enough that the benefit of parallelization outweighs the overhead
  - In the following example, what happens when we parallelize `isPrime`?

19 / 24

## Overhead of Fork/Join

- Forking and joining creates additional overhead
  - wrap the computation in an object
  - submit object to a queue of tasks
  - workers go through the queue to execute tasks



18 / 24

## Fork/Join in Parallel Streams

```
public static boolean isPrime(int n) {
    return IntStream.range(2, (int) Math.sqrt(n) + 1)
        .noneMatch(x -> n % x == 0);
}

public static void main(String[] args) {
    if (args.length != 0) {
        System.setProperty(
            "java.util.concurrent.ForkJoinPool.common.parallelism",
            args[0]);
    }
    System.out.println("Number of worker threads: " +
        ForkJoinPool.commonPool().getParallelism());

    Instant start = Instant.now();
    long howMany = IntStream.range(2_000_000, 3_000_000)
        .parallel()
        .filter(x -> isPrime(x))
        .count();
    Instant stop = Instant.now();

    System.out.println(howMany + " : " +
        Duration.between(start, stop).toMillis() + "ms");
}
```

20 / 24

## Comparing Sequential and Parallel Streams

□ Suppose given the following task unit

```
class OneSecondTask {
    int ID;

    public OneSecondTask(int ID) {
        this.ID = ID;
    }

    public int compute() {
        System.out.println(Thread.currentThread().getName());
        try {
            Thread.sleep(1000);
        } catch (InterruptedException e) {
            throw new RuntimeException(e);
        }
        return ID;
    }
}
```

## Comparing Sequential and Parallel Streams

```
public static void parallelStreamRun(List<OneSecondTask> tasks) {
    Instant start = Instant.now();
    List<Integer> result = tasks.parallelStream()
        .map(x -> x.compute())
        .collect(Collectors.toList());
    Instant stop = Instant.now();
    System.out.print(result + " ");
    System.out.println(Duration.between(start,stop).toMillis() + "ms")
}
```

□ Parallel stream on 4 worker threads:

```
main
ForkJoinPool.commonPool-worker-1
ForkJoinPool.commonPool-worker-3
ForkJoinPool.commonPool-worker-2
ForkJoinPool.commonPool-worker-4
ForkJoinPool.commonPool-worker-3
ForkJoinPool.commonPool-worker-2
ForkJoinPool.commonPool-worker-4
ForkJoinPool.commonPool-worker-1
ForkJoinPool.commonPool-worker-3
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] 3006ms
```

## Comparing Sequential and Parallel Streams

```
public static void sequentialRun(List<OneSecondTask> tasks) {
    Instant start = Instant.now();
    List<Integer> result = tasks.stream()
        .map(x -> x.compute())
        .collect(Collectors.toList());
    Instant stop = Instant.now();
    System.out.print(result + " ");
    System.out.println(Duration.between(start,stop).toMillis() + "ms")
}
```

□ Sequential stream on 4 worker threads:

```
main
main
main
main
main
main
main
main
main
main
main
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] 10003ms
```

## Lecture Summary

- Familiarity with the use of parallel streams
- Adherence to rules for parallelizing streams
- Appreciate fork and join
  - Thread pools
  - Fork/join overhead