

CS2030 Lecture 11


Fork/Join Framework

Henry Chia (hchia@comp.nus.edu.sg)

Semester 1 2018 / 2019

Fork and Join

- Given the following program fragment and *computation* graph



- `f(a)` invoked before `g(b)` and `h(b)`; `n(c,d)` invoked after
- How about the order of `g(b)` and `h(b)`?
 - If `g` and `h` does not produce side effects, then parallelize
 - **Fork** task `g` to execute at the same time as `h`, and **join** back task `g` later

1 / 22

3 / 22

Lecture Outline

- Fork and join tasks
- Java's Fork/join framework
 - Sub-classing a `RecursiveTask`
- Thread pools
 - Global queue
 - Local deque (double-ended queue)
- Work stealing
- Order of fork and join
- Overhead of fork and join
- Fork/join in parallel streams

Example: Summing an Array... *Recursively*

```
class Summer {
    static int threshold;

    static int sumLeftRight(int[] array, int low, int high) {
        if (high - low < threshold) {
            int sum = 0;
            for (int i = low; i <= high; i++) {
                sum += array[i];
            }
            return sum;
        } else {
            int middle = (low + high) / 2;
            int leftSum = sumLeftRight(array, low, middle);
            int rightSum = sumLeftRight(array, middle + 1, high);
            return leftSum + rightSum;
        }
    }
}

int[] array = IntStream.rangeClosed(1, 10).toArray();
Summer.threshold = Integer.parseInt(args[0]);
int sum = sumLeftRight(array, 0, array.length - 1);
```

2 / 22

4 / 22


Transform to a Recursive Class

```
class Summer {
    int[] array;
    int low;
    int high;
    static int threshold;

    Summer(int[] array, int low, int high) {
        this.array = array;
        this.low = low;
        this.high = high;
    }

    int compute() {
        if (high - low < threshold) {
            int sum = 0;
            for (int i = low; i <= high; i++) {
                sum += array[i];
            }
            return sum;
        } else {
            int middle = (low + high) / 2;
            Summer left = new Summer(array, low, middle);
            Summer right = new Summer(array, middle + 1, high);
            return left.compute() + right.compute();
        }
    }
}
```

Thread Pools



5 / 22

7 / 22

Subclassing RecursiveTask<T> for Fork/Join

```
import java.util.concurrent.RecursiveTask;

class Summer extends RecursiveTask<Integer> {
    ...
    @Override
    protected Integer compute() {
        if (high - low < threshold) {
            int sum = 0;
            for (int i = low; i < high; i++) {
                sum += array[i];
            }
            return sum;
        } else {
            int middle = (low + high) / 2;
            Summer left = new Summer(low, middle, array);
            Summer right = new Summer(middle, high, array);

            left.fork();
            return right.compute() + left.join();
        }
    }
}
```

Thread Pools

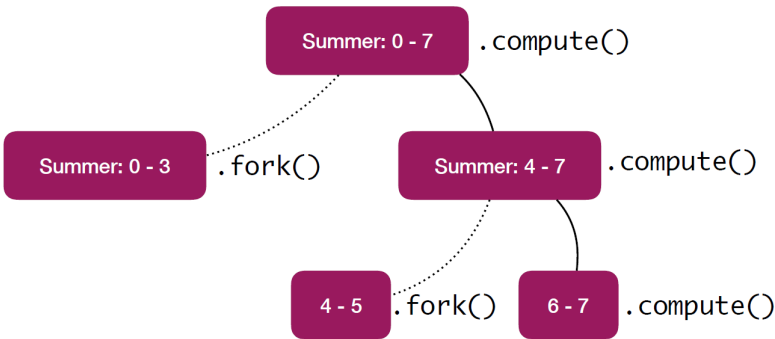

- Java maintains a pool of *worker threads*
 - Each thread is an abstraction of a running task
 - Task submitted to the pool for execution, and joins the global queue or worker queue
 - Worker thread picks a task from the queue to execute
- ForkJoinPool is the class that implements the thread pool for RecursiveTask (a sub-class of ForkJoinTask)
- To submit task to the thread pool for execution, either
 - task.compute() that invokes task immediately; may result in stack overflow if too many recursive tasks
 - invoke(task) that gets the task to join the queue, waiting to be carried out by a worker (recommended)

6 / 22

8 / 22


Example: Summing an Array

- Summing array of eight elements with threshold set to 2



Queuing of Forked Tasks


- main thread performs compute on task [0,7]



#1

#2

- main forks task [0,3] to local queue, then computes [4,7]



0-3

#1

#2

Example: Summing an Array

```
@Override
protected Integer compute() {
    System.out.println(low + "," + high + ":" + Thread.currentThread().getName());
    if (high - low < threshold) {
        int sum = 0;
        for (int i = low; i < high; i++) {
            sum += array[i];
        }
        return sum;
    } else {
        int middle = (low + high) / 2;
        Summer left = new Summer(array, low, middle);
        Summer right = new Summer(array, middle + 1, high);
        left.fork();
        return right.compute() + left.join();
    }
}
```


- Running with ForkJoinPool.common.parallelism=2

0,7:main
4,7:main
6,7:main
0,3:ForkJoinPool.commonPool-worker-1
4,5:ForkJoinPool.commonPool-worker-2
2,3:ForkJoinPool.commonPool-worker-1
0,1:ForkJoinPool.commonPool-worker-2

0,7:main
4,7:main
6,7:main
4,5:main
0,3:ForkJoinPool.commonPool-worker-2
0,1:ForkJoinPool.commonPool-worker-1
2,3:ForkJoinPool.commonPool-worker-2

Work Stealing


- Worker #1 steals task [0-3]



#1

#2

- main forks [4-5] (worker #2 steals); main computes [6-7]




#1


#2

Work Stealing

- Worker #1 forks [0-1] and computes [2-3]




- Worker #2 completes [4-5] returns, and steals [0-1]




13 / 22

Another Possible Scenario...

- Worker #2 steals task [0-3]




- Worker #2 forks [0-1] (worker #1 steals) and computes [2-3]



15 / 22

Local Deque (Double-Ended Queue)

- main forks task [0-3] to local queue
- Before any work stealing, main computes [4-7] which forks [4-5] and computes [6-7]




- Local task queue is a double ended queue
 - Forked tasks added to the **head** of the queue
 - Steal tasks from the **end of the queue**
 - Rational: bigger tasks are stolen; smaller ones self-served

14 / 22

Another Possible Scenario...

- Suppose main completes [6-7], but worker #1 has not
- Worker #2 completes [2-3], but is now blocked waiting for worker #1 to return with a value
- main can service 4-5 from the head of its queue



- Indeed, if there are no workers, main completes everything on his own

16 / 22

Order of Fork/Join

```
@Override
protected Integer compute() {
    System.out.println(low + "," + high + ":" + Thread.currentThread().getName());
    if (high - low < threshold) {
        int sum = 0;
        for (int i = low; i < high; i++) {
            sum += array[i];
        }
        return sum;
    } else {
        int middle = (low + high) / 2;
        Summer left = new Summer(array, low, middle);
        Summer right = new Summer(array, middle + 1, high);

        left.fork();
        right.fork();


        return right.compute() + left.compute();
    }
}
```

- How about using only forks and joins
- Does the ordering matter?

17 / 22

Overhead of Fork/Join

- Forking and joining creates additional overhead
 - wrap the computation in an object
 - submit object to a queue of tasks
 - workers go through the queue to execute tasks



19 / 22

Order of Fork/Join

- Fork-join pair acts like a call (fork) and return (join) from a parallel recursive function.
 - Returns (joins) should be performed innermost-first
 - `a.fork(); b.fork(); b.join(); a.join();` is likely to be substantially more efficient than joining a before b.
- Work-stealing threadpools have a fixed number of threads
 - Any blocking operation in one of these threads will reduce overall performance
- Will all threads be blocked?
 - Synchronization primitive classes recognize when a thread is about to be blocked and create a *compensation thread* before the current thread blocks

18 / 22

Fork/Join in Parallel Streams

- `parallel()` runs fork to create sub-tasks running the same chain of operations on sub-streams
- `combiner` in `reduce` runs join to combine the results
- Parallelizing a trivial task actually creates more work in terms of parallelizing overhead

```
IntStream.range(2, (int) Math.sqrt(n) + 1)
    .parallel()
    .noneMatch(x -> n % x == 0);
```
- Parallelization is worthwhile if the task is complex enough that the benefit of parallelization outweighs the overhead
- In the following example, what happens when we parallelize `isPrime`?

20 / 22

Fork/Join in Parallel Streams

```
public static boolean isPrime(int n) {  
    return IntStream.range(2, (int) Math.sqrt(n) + 1)  
        .noneMatch(x -> n % x == 0);  
}  
  
public static void main(String[] args) {  
    System.out.println("Number of worker threads: " +  
        ForkJoinPool.commonPool().getParallelism());  
  
    Instant start = Instant.now();  
    long howMany = IntStream.range(2_000_000, 3_000_000)  
        .parallel()  
        .filter(x -> isPrime(x))  
        .count();  
    Instant stop = Instant.now();  
  
    System.out.println(howMany + " : " +  
        Duration.between(start, stop).toMillis() + "ms");  
}
```

21 / 22

Lecture Summary

- Appreciate the use of fork and join in parallel/concurrent programming
- Understand how tasks are forked in a local deque, and why the ordering of forks and joins matter
- Understand how work stealing distributes tasks among worker threads
- Appreciate the overhead involved in parallelizing using fork and join

22 / 22